

UNCLASSIFIED

DECENTRALIZING THE PROCESSING UNIT: DISTRIBUTING COMPUTE THROUGHOUT THE SYSTEM

DAVID DONOFRIO, CHIEF HARDWARE ARCHITECT

JOHN D. LEIDEL, CHIEF SCIENTIST

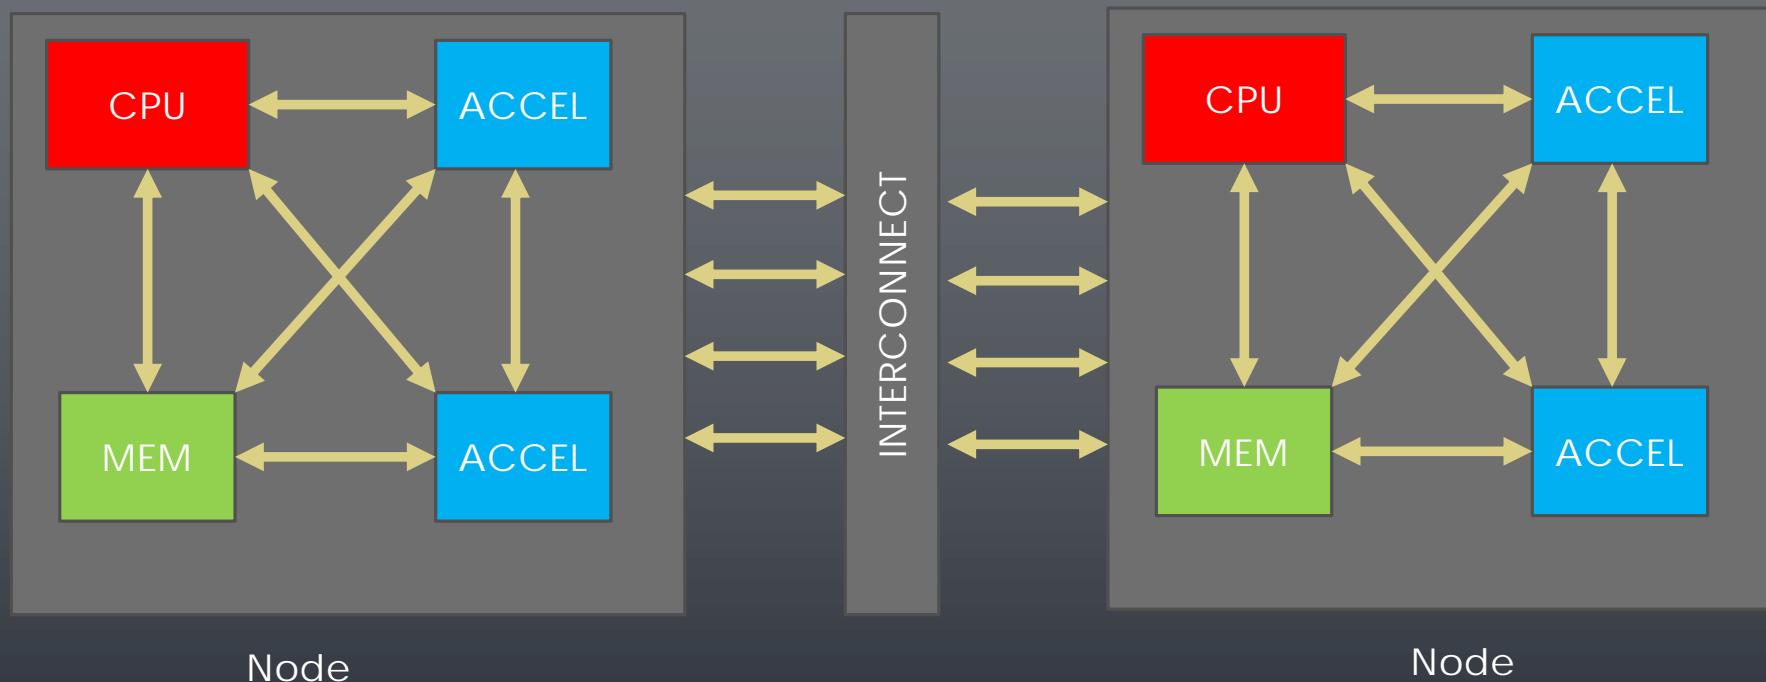
{DDONOFRIO,JLEIDEL}@TACTCOMPLABS.COM

2020/28/10

UNCLASSIFIED

UNCLASSIFIED

SYSTEMS TODAY



UNCLASSIFIED

UNCLASSIFIED



SPECIALIZATION AS THE PATH TO PERFORMANCE

- PROCESS SCALING TAPERING OFF
- SPECIALIZED ARCHITECTURES ARE EVERYWHERE
 - GPUS
 - AI ACCELERATORS
 - RE-CONFIGURABLE
 - MATRIX ACCELERATORS
- INCREASING HETEROGENEITY IN SYSTEMS
 - HETEROGENOUS IN BOTH TYPES OF COMPUTATION AS WELL AS WHERE COMPUTING IS PERFORMED
 - AS COMPUTING BECOMES MORE DECENTRALIZED WE WILL NEED A GREATER VARIETY OF ACCELERATORS
 - THIS DRIVES A NEED FOR FRAMEWORKS THAT ALLOW RAPID DESIGN SPACE EXPLORATION

UNCLASSIFIED

UNCLASSIFIED



A BIT OF HISTORY

- LBNL & TEXAS TECH WERE FUNDED TO DEVELOP A SPECIAL-PURPOSE, RISC-V BASED SOC
 - FPGA-BASED
 - HMC MEMORIES, CUSTOM ISA EXTENSIONS, SCRATCHPADS
- PHD DISSERTATION
 - GOBLINCORE-64: RISC-V BASED SOC FOR DATA INTENSIVE COMPUTING
 - ISA EXTENSIONS, PIPELINE CHANGES, SMT
 - NO CACHES, MEMORY COALESCING



UNCLASSIFIED

UNCLASSIFIED



SYSTEM ARCHITECT GEN1

- THIS “RINSE AND REPEAT” DESIGN FLOW CAN BE AUTOMATED!
- SYSTEM ARCHITECT GEN1 WAS BORN AS A TOOLCHAIN FOR DEVELOPING RISC-V EXTENSIONS
- UTILIZED COMPILER-TECHNIQUES TO HANDLE DEPENDENCE ANALYSIS OF DESIGN PARAMETERS TO RISC-V INTERNALS
- **INPUT:** RISC-V CORE & DESIGN PARAMETERS
- **OUTPUT:** CHISEL RTL FOR ENTIRE SOC & LLVM COMPILER



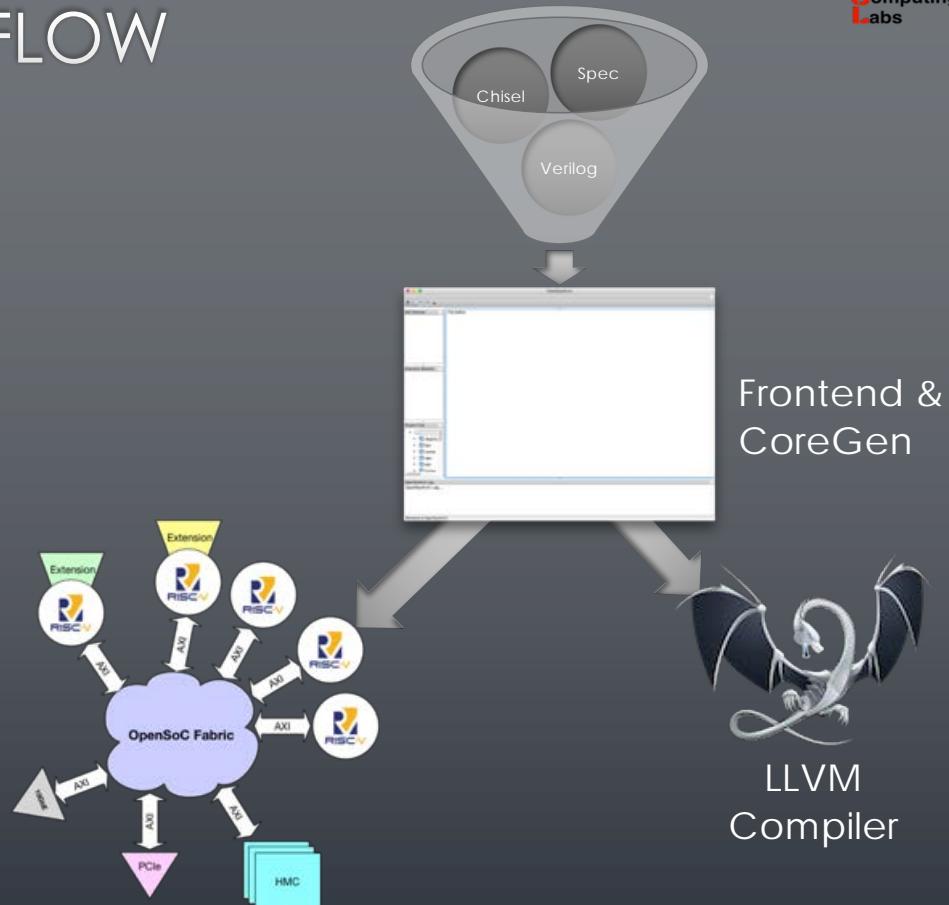
UNCLASSIFIED

UNCLASSIFIED



SYSARCH TOOL CHAIN & FLOW

- OPENSoC Tool Chain & Flow consists of several integral moving parts:
 - OPENSoC System Architect Frontend
 - COREGEN IR & Backend
 - OPENSoC Fabric NoC Interconnect
 - RISC-V Standard Cores (in Chisel)
 - LLVM Compiler Infrastructure
- Users input design specs and extensions via Frontend
- COREGEN ensures design correctness and continuity
 - Also generates backend Chisel and LLVM implementation
- OPENSoC Fabric “glues” all intermediate modules together with Scalable NoC



UNCLASSIFIED

UNCLASSIFIED



SYSTEM ARCHITECT GEN2

- SEVERAL EARLY USERS POSED AN INTERESTING QUESTION:
- ***WHAT ABOUT DEVELOPING ARBITRARY ISA's?***
- ***GEN2 IS A COMPLETE RE-WRITE OF THE INFRASTRUCTURE TO SUPPORT ARBITRARY ISA DEVELOPMENT***
 - NEW INTERMEDIATE REPRESENTATION (FORMATTED IN YAML)
 - NEW DEPENDENCE ANALYSIS ENGINE (COREGEN)
 - HIGH LEVEL INSTRUCTION DEFINITION LANGUAGE (STONECUTTER)
 - SUPPORT FOR PLUGINS



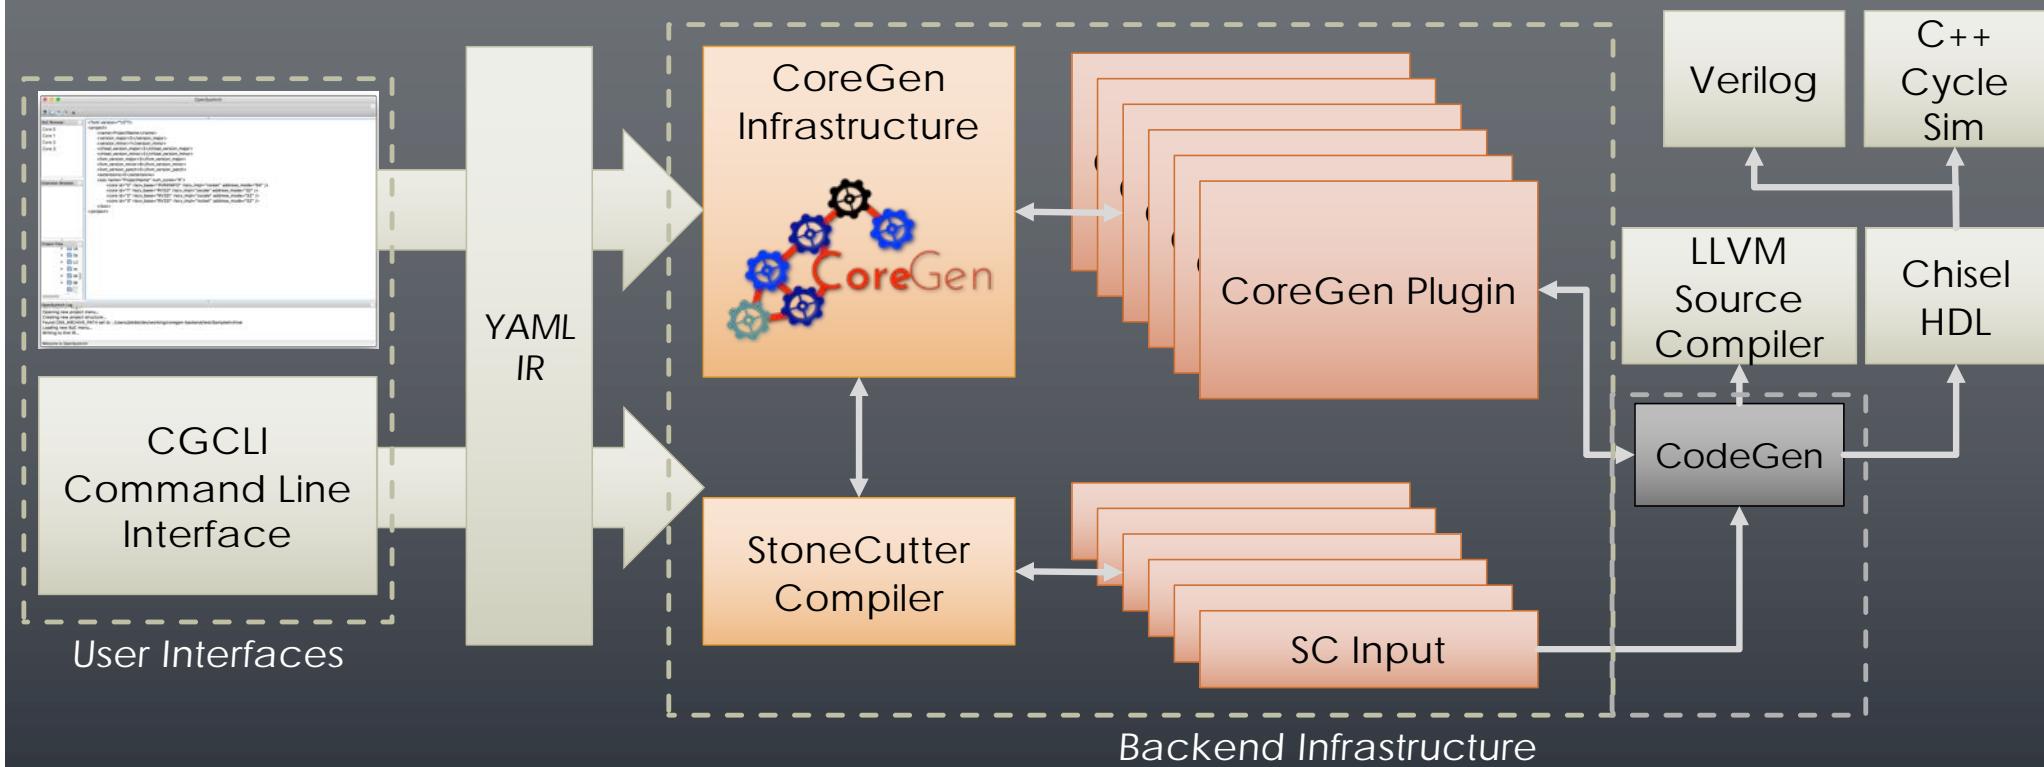
Maintain verification performance similar to a modern optimizing compiler!

UNCLASSIFIED

UNCLASSIFIED



GEN2 HIGH LEVEL DESIGN

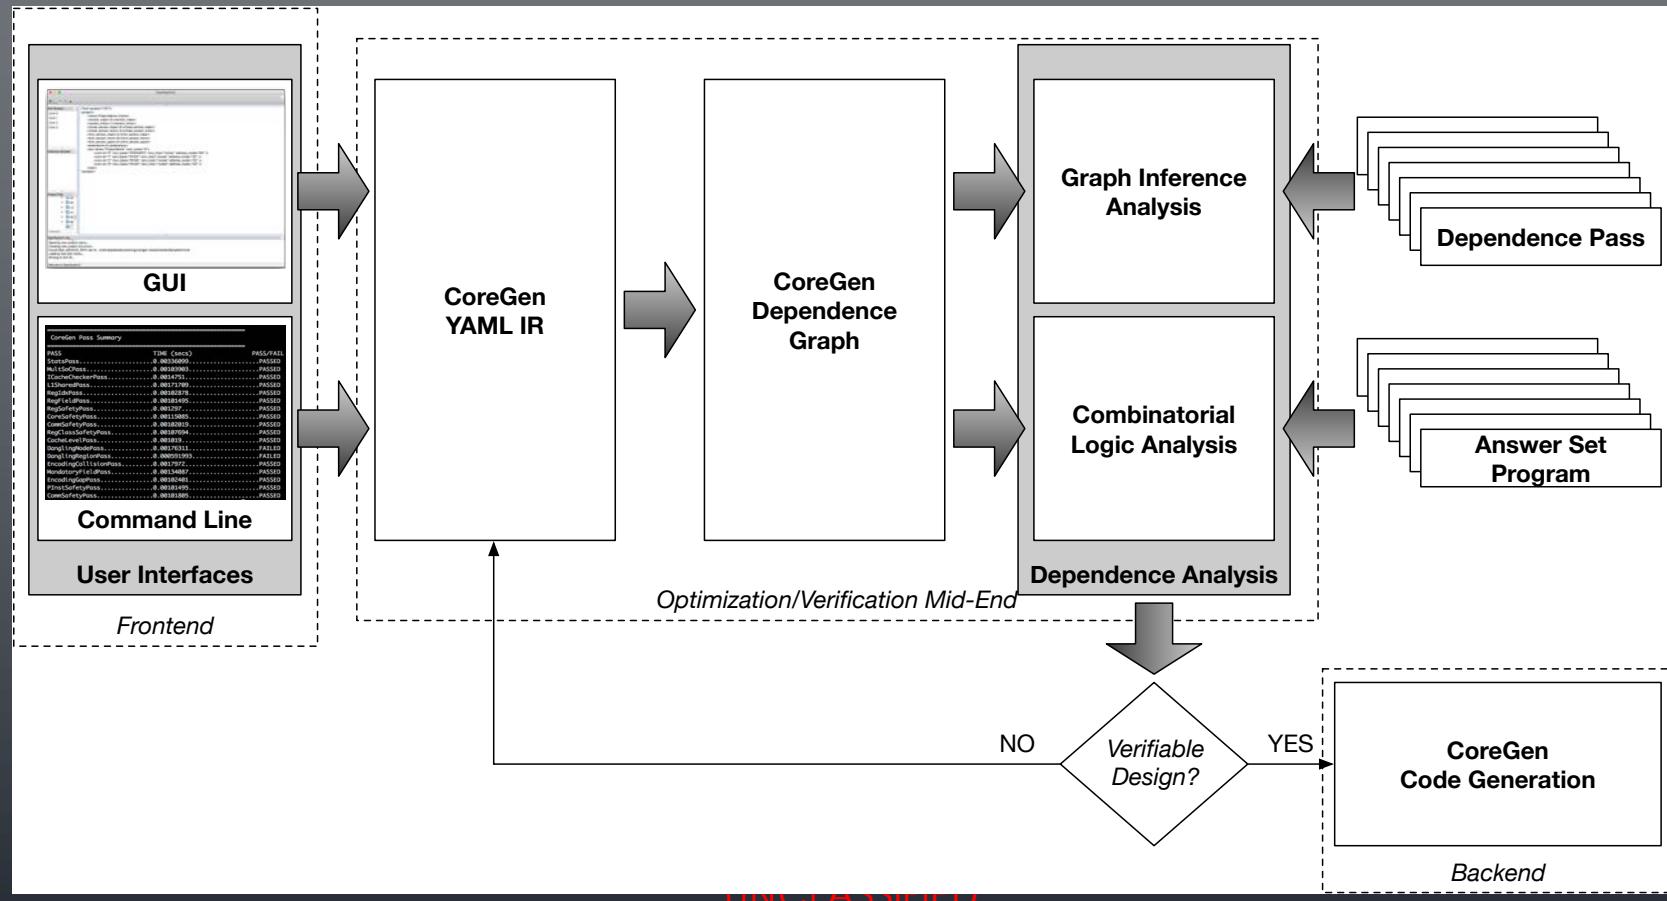


UNCLASSIFIED

UNCLASSIFIED



COREGEN VERIFICATION INFRASTRUCTURE

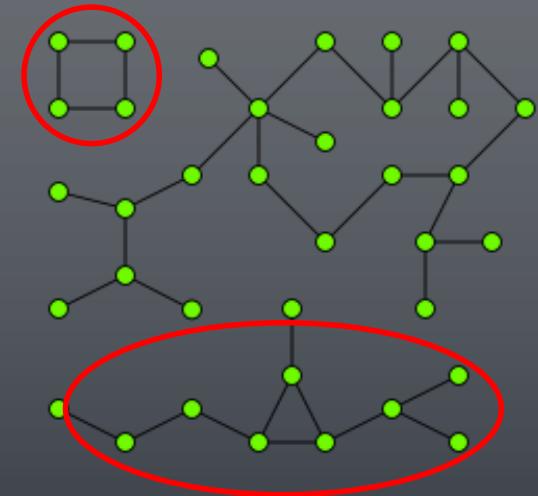


UNCLASSIFIED



SYSTEM ARCHITECT GEN2: RAPID DESIGN GENERATION AND ANALYSIS

- GRAPH-BASED DEPENDENCE ANALYSIS FOR DESIGN VERIFICATION
- HARDWARE MODULES ARE DEFINED AS NODES IN THE DAG INSIDE COREGEN
 - DEPENDENCE GRAPH PROCESSED IN MULTIPLE STAGES TO EXPOSE CONNECTIVITY, COMMUNICATION LINKS, INSTRUCTION AND REGISTER ENCODINGS
- MULTIPLE PASSES PERFORMED ON THE DAG FOR DESIGN VERIFICATION
 - ELIMINATION OF DANGLING NODES / REGIONS
 - EARLY DETECTION TO AVOID ISSUES WITH LATER SYNTHESIS OF THE DESIGN
 - ISA ENCODING COLLISION CHECK
 - EARLY DETECTION WITH DECODE LOGIC AND ASSEMBLY/DISASSEMBLY
 - COMBINATORIAL LOGIC ANALYSIS
 - USER-DEFINED PASSES COMING SOON



[https://en.wikipedia.org/wiki/Component_\(graph_theory\)](https://en.wikipedia.org/wiki/Component_(graph_theory))

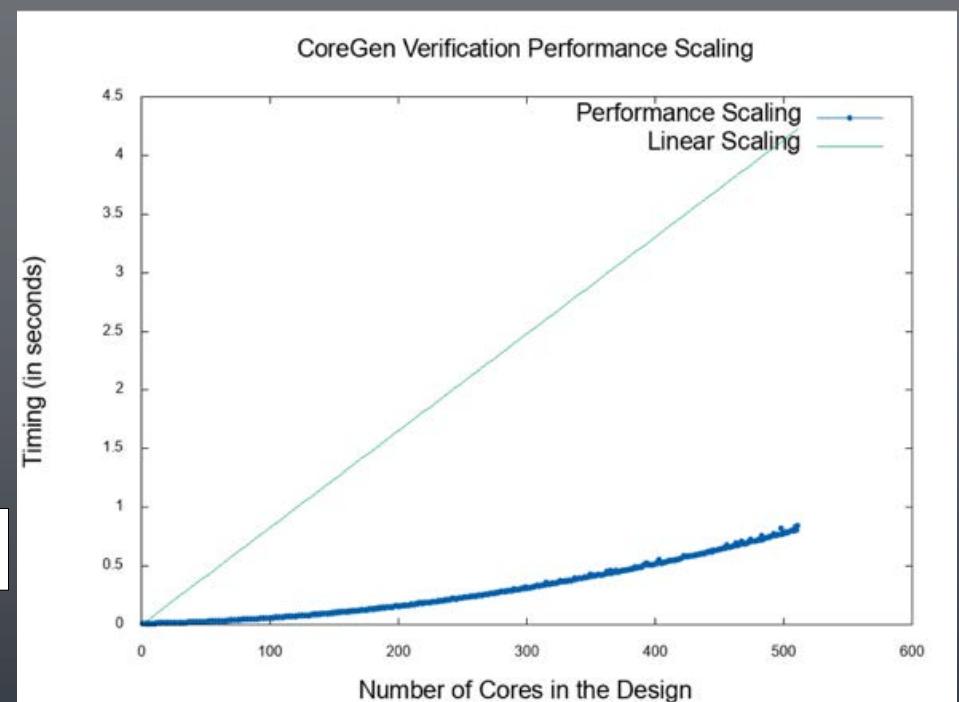
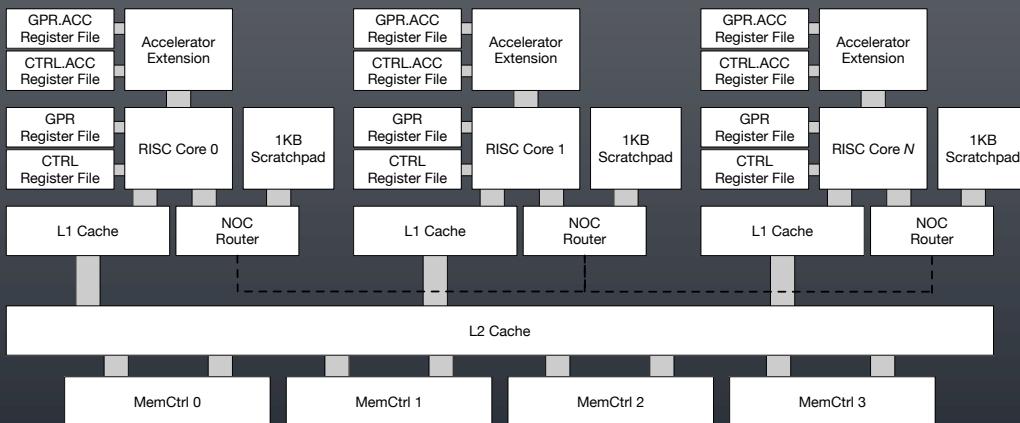
UNCLASSIFIED

UNCLASSIFIED



SAMPLE PERFORMANCE

- HETEROGENEOUS, TILED DESIGN WITH MULTI-LEVEL CACHES AND NOC
 - MULTIPLE ISA's WITH RISC CORES AND ATTACHED ACCELERATORS
- DESIGN SCALED FROM 1 TO 511 TILES
- **FULL VERIFICATION RUN OF 511 CORES < 1 SEC**



UNCLASSIFIED

UNCLASSIFIED



PROGRAMMABLE ACCELERATOR NETWORKS: AMORTIZING HPC INTERCONNECTS

UNCLASSIFIED

UNCLASSIFIED



PAN: PROGRAMMABLE ACCELERATOR NETWORK

- CURRENT HPC INTERCONNECTS MOVE SIGNIFICANT DATA – AND CONSUME SIGNIFICANT POWER
 - OPPORTUNITY TO AMORTIZE THE POWER AND LATENCY INCURRED IN TRAVERSING THE NETWORK BY COMPUTE EMBEDDED IN THE INTERCONNECT
- PAN COMBINES SDN CONCEPTS WITH HPC INTERCONNECTS EMBEDDED WITH IN-SITU COMPUTE ACCELERATORS
- GOAL IS TO SUPPORT BOTH TRADITIONAL NETWORK OPERATIONS AS WELL AS COMPUTE-IN-NETWORK
- ACCELERATE APPLICATIONS WITH LOW COMPUTE EFFICIENCY AND SMALL MESSAGES BY COMBINING THESE OPERATIONS INTO A SINGLE SOLUTION
 - LARGE SCALE SPARSE SOLVERS AND GRAPH ALGORITHMS.

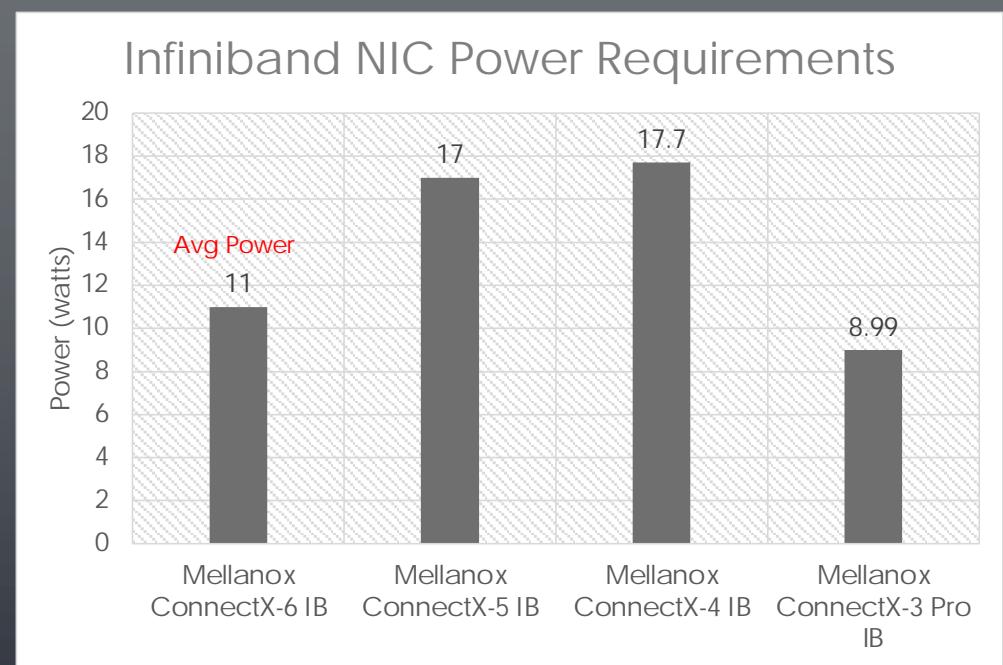
UNCLASSIFIED

UNCLASSIFIED



INFINIBAND POWER REQUIREMENTS

- MELLANOX INFINIBAND NIC POWER REQUIREMENTS
- DATA WAS OBTAINED FROM MELLANOX DOCUMENTATION
- ** CONNECTX-6 IS AVERAGE POWER
- ** CONNECTX-{3-5} IS PEAK POWER

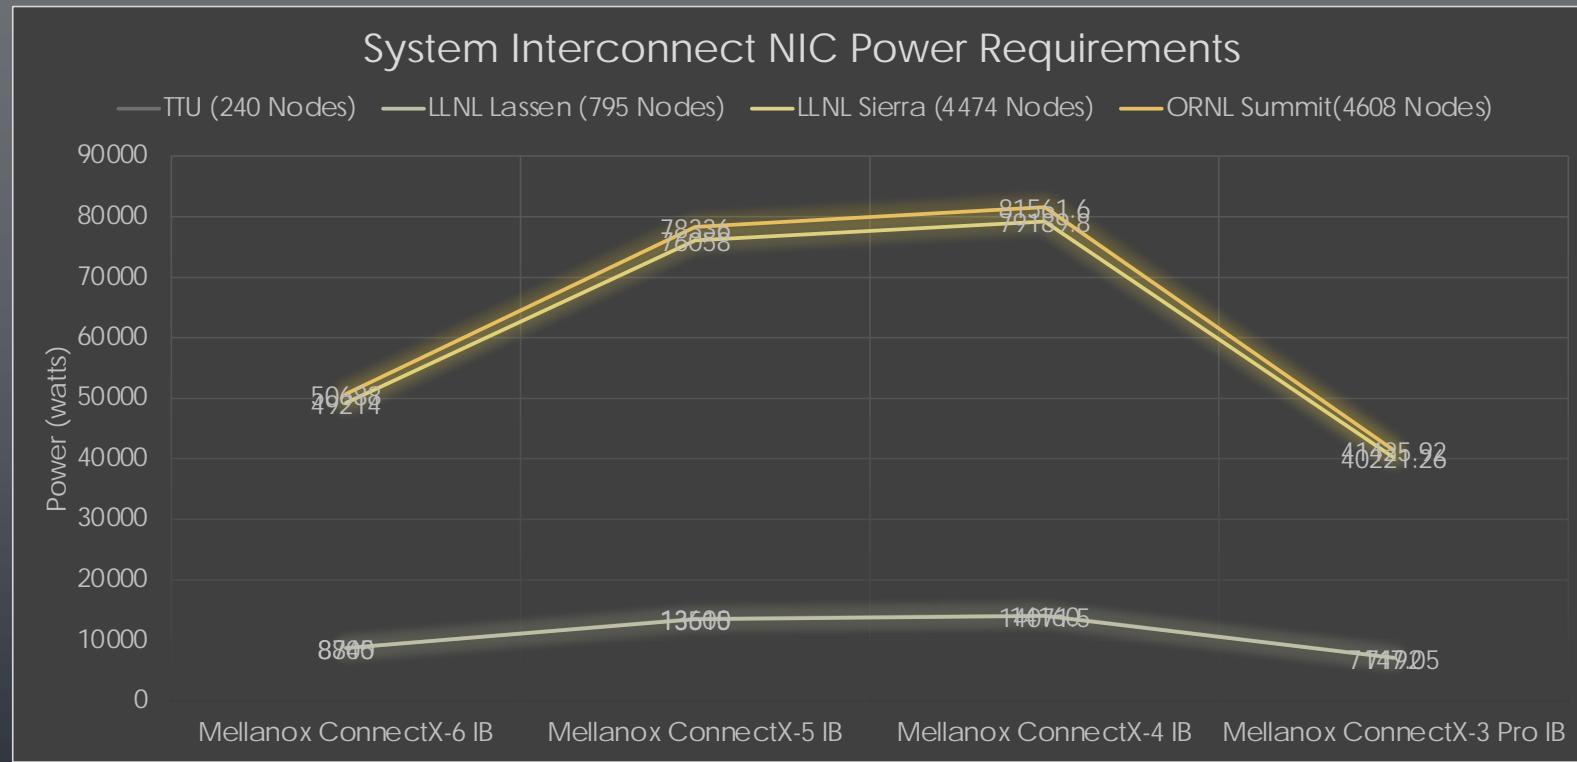


UNCLASSIFIED

UNCLASSIFIED



SCALABLE SYSTEM INTERCONNECT POWER



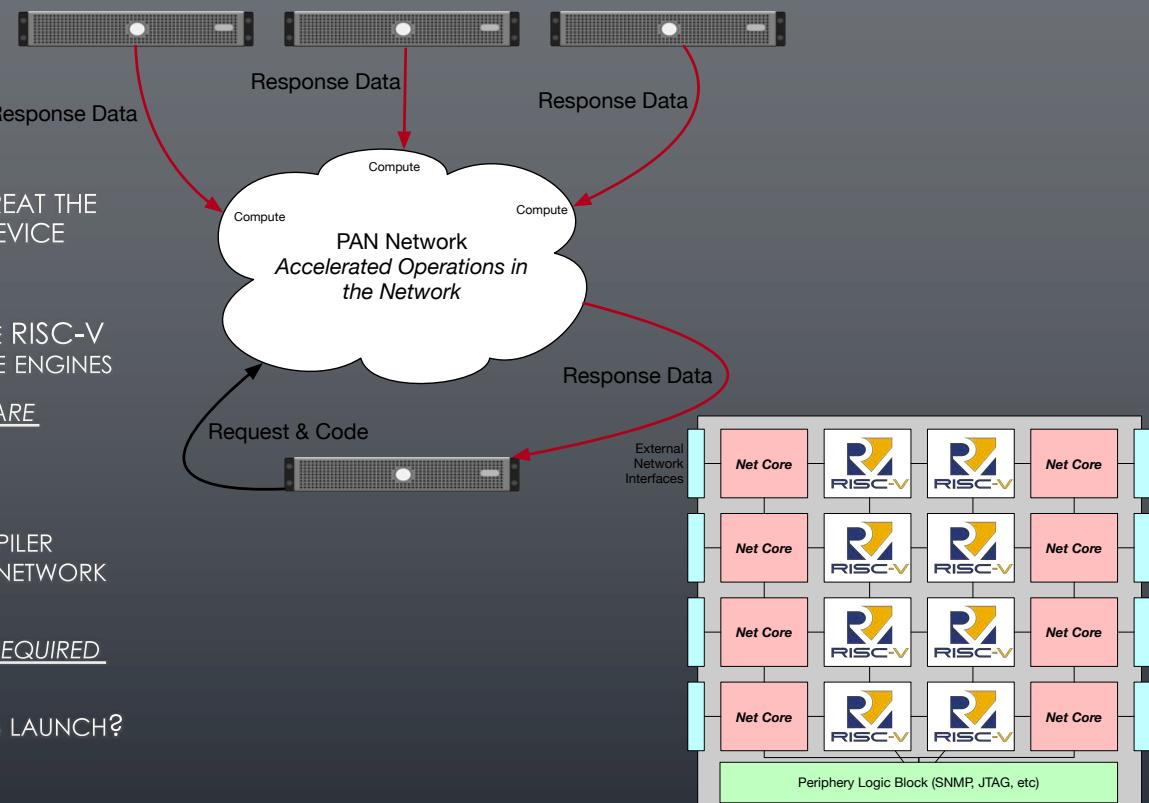
UNCLASSIFIED

UNCLASSIFIED



PROGRAMMABLE ACCELERATOR NETWORKS

- DEVELOP SIMULATION AND DEMONSTRATION HARDWARE/SOFTWARE IP TO PERMIT USERS TO TREAT THE INTERCONNECT AS A **COMPUTE** ACCELERATOR DEVICE
- HARDWARE IP:
 - COUPLE NETWORK BLOCKS TO ONE OR MORE RISC-V DEVICES TO SERVE AS THE NETWORK COMPUTE ENGINES
 - WHAT ADDITIONAL ISA/UARCH EXTENSIONS ARE REQUIRED?
- SOFTWARE IP:
 - DEVELOP PROGRAMMING MODEL AND COMPILER EXTENSIONS TO PERMIT SEAMLESS ACCESS TO NETWORK COMPUTE CORES
 - WHAT ADDITIONAL UARCH EXTENSIONS ARE REQUIRED FOR SEAMLESS FUNCTIONALITY?
 - DEBUGGING? CONTEXT SAVE/RESTORE? JOB LAUNCH?



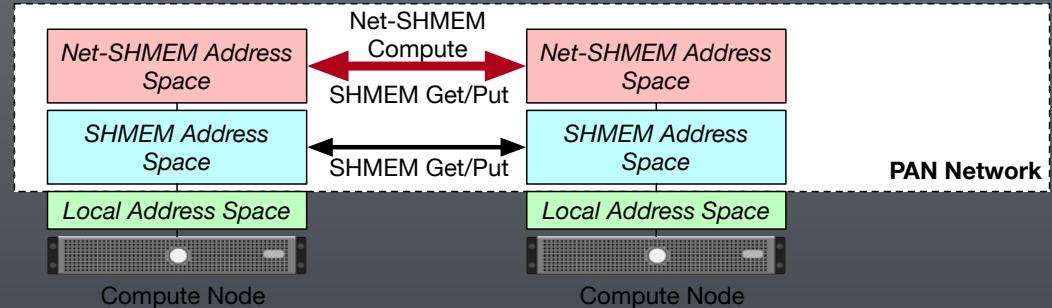
UNCLASSIFIED

UNCLASSIFIED



EXAMPLE PROGRAMMING MODEL: NET-SHMEM

- EXTENDS OPENSHMEM PROGRAMMING MODEL TO EXPOSE NETWORK MEMORY/COMPUTE SPACE
- THE NIC/SWITCH MEMORY IS EXPOSED VIA A SEPARATE NAMED ADDRESS SPACE
- UTILIZES RUDIMENTARY SHMEM OPERATIONS FOR DATA MOTION BETWEEN HOST-NETWORK SPACE AND BETWEEN NETWORK-NETWORK SPACE



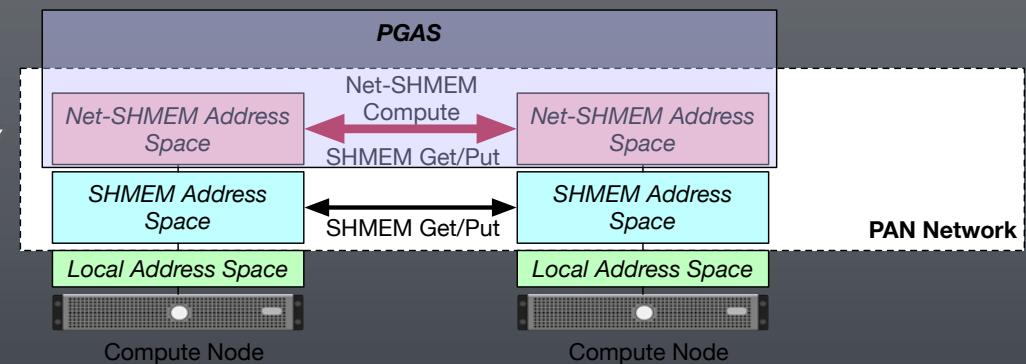
UNCLASSIFIED

UNCLASSIFIED



EXAMPLE PROGRAMMING MODEL: PGAS

- EXTENDS OPENSHMEM, CHAPEL, UPC/UPC++ TO PROMOTE NETWORK MEMORY SPACE AS DIRECTLY ACCESSIBLE PGAS MEMORY
- DMA OPERATIONS FOR DATA MOTION BETWEEN HOST/NETWORK MEMORY SPACE
- EXTENDED LAUNCH MECHANISMS FOR DIRECT ACCESS TO NETWORK COMPUTE FACILITIES



UNCLASSIFIED

UNCLASSIFIED



PAN SIMULATION INFRASTRUCTURE

UNCLASSIFIED

UNCLASSIFIED



REV SIMULATION INFRASTRUCTURE

- NATIVE SST COMPONENT THAT EXISTS OUTSIDE THE NORMAL SST TREE
 - DOES NOT REQUIRE ANY MODIFICATIONS TO SST
- TOP-LEVEL COMPONENT IS MAPPED TO A CPU DEVICE
- THE CPU (OR CPUS) IMPLEMENT SINGLE OR MULTI-CORE RISC-V DEVICES
- EACH “SOCKET” CAN CONTAIN HETEROGENEOUS CORES
- UNLIKE OTHER CPU MODELS, REV CONTAINS A **REAL** BINARY LOADER
 - REV RUNS APPLICATIONS! NOT JUST SIMPLE TESTS
 - IN-ORDER, CYCLE-BASED MODEL WITH CONFIGURABLE MEMORY LATENCY AND INSTRUCTION LATENCY
- CONTAINS A SIMPLE MERLIN NIC (RevNIC)

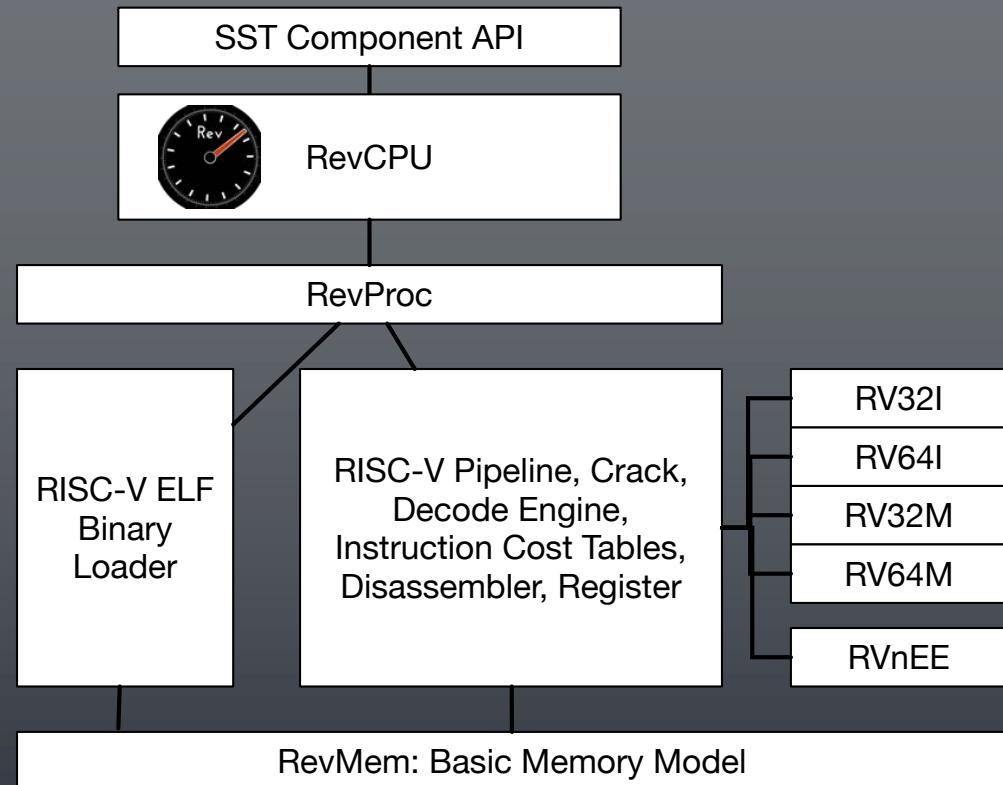
UNCLASSIFIED

UNCLASSIFIED



REV ARCHITECTURE

- CYCLE BASED CPU MODEL
 - CURRENTLY IN-ORDER ONLY
 - INSTRUCTION COST TABLES ARE CONFIGURABLE AT RUNTIME (TEXT FILE)
- EXECUTES NATIVE ELF BINARIES
 - NO CUSTOM COMPILER REQUIRED!
- SUPPORT FOR NEW INSTRUCTION EXTENSIONS WITH MINIMAL EFFORT
- CURRENTLY ALL INTEGER INSTRUCTIONS ARE COMPLETE
 - WORKING ON FLOATING POINT
- RELEASED UNDER APACHE 2.0 LICENSE
- [HTTPS://GITHUB.COM/TACTCOMPLABS/REV](https://github.com/tactcomplabs/rev)

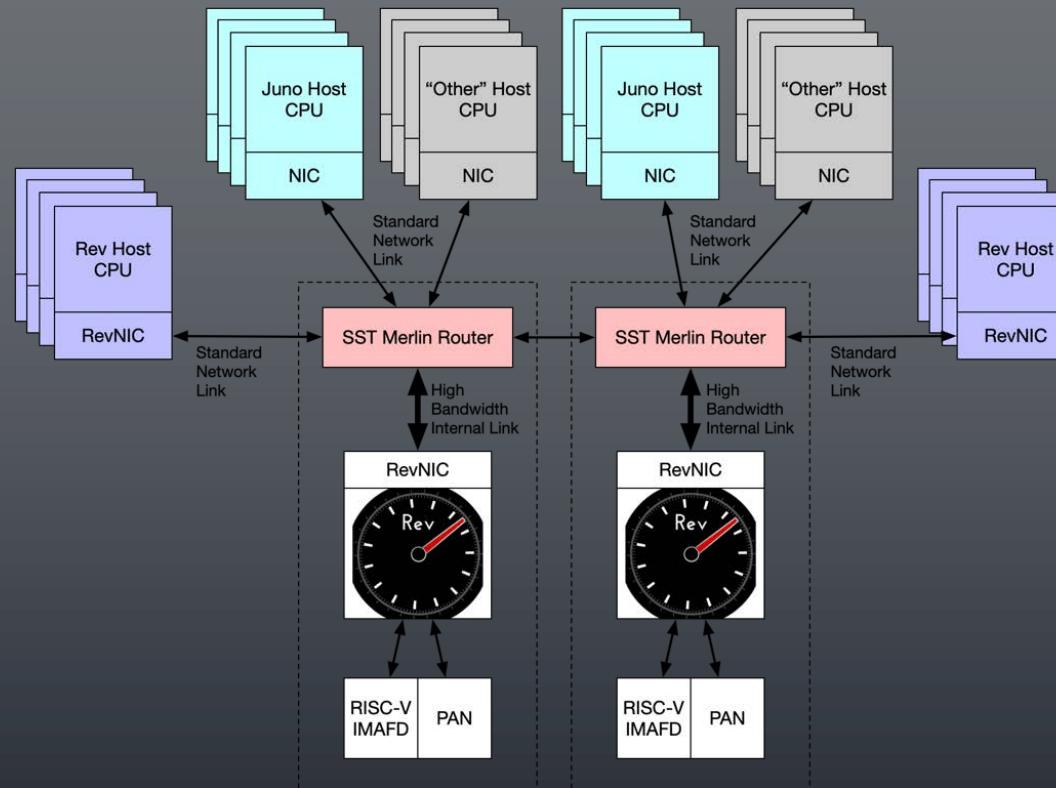


UNCLASSIFIED

UNCLASSIFIED



REV SIMULATION INFRASTRUCTURE



UNCLASSIFIED

UNCLASSIFIED



PAN SIMULATION INFRASTRUCTURE

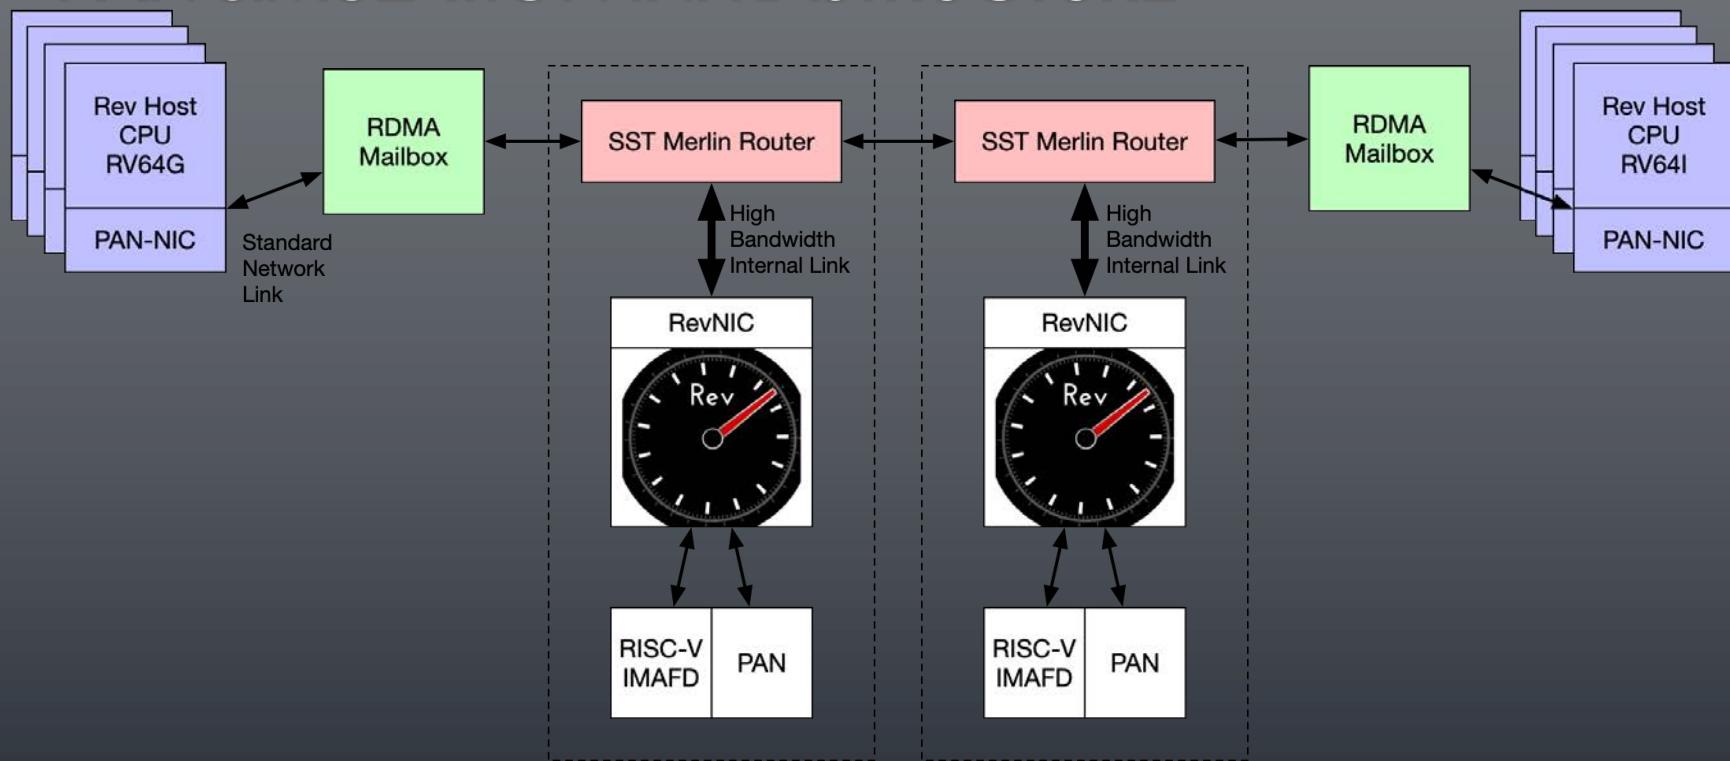
- PAN EXTENDS THE BASE REV CPU COMPONENT WITH ADDITIONAL SUB-COMPONENTS
- PAN-NET (PAN-NIC) IMPLEMENTS THE PAN COMMAND PACKET INFRASTRUCTURE
 - THIS INCLUDES AN RDMA MAILBOX COMPONENT IMPLEMENTED AT THE CPU LEVEL
 - THE PAN-NIC DOES NOT REQUIRE THE PAN RISC-V ISA EXTENSIONS (VANILLA HOST, ACCELERATED NETWORK)
- PAN RISC-V EXTENSION: IMPLEMENTS "FUTURE" OPERATIONS
 - THE REV INFRASTRUCTURE STILL SUPPORTS HETEROGENEOUS OPERATIONS

UNCLASSIFIED

UNCLASSIFIED



PAN SIMULATION INFRASTRUCTURE



UNCLASSIFIED

UNCLASSIFIED



PAN COMMAND SET

- SIMULATOR SUPPORTS ALL THE PACKET COMMAND TYPES
- SYNCHRONOUS/ASYNCHRONOUS GETS/PUTS
- JOB MANAGEMENT
- DEBUGGING

Packet	Packet Type	Opc Encoding
Synchronous Get	<i>Base</i>	0b00000000
Synchronous Put	<i>Base</i>	0b00000100
Asynchronous Get	<i>Base</i>	0b00001000
Asynchronous Put	<i>Base</i>	0b00001100
Synchronous Streaming Get	<i>Streaming</i>	0b0001
Synchronous Streaming Put	<i>Streaming</i>	0b0101
Asynchronous Streaming Get	<i>Streaming</i>	0b1001
Asynchronous Streaming Put	<i>Streaming</i>	0b1101
Execute	<i>Base</i>	0b00010000
Status	<i>Base</i>	0b00100000
Cancel	<i>Base</i>	0b00110000
Reserve	<i>Base</i>	0b01000000
Revoke	<i>Base</i>	0b01010000
Halt	<i>Base</i>	0b01100000
Resume	<i>Base</i>	0b01110000
Read Register	<i>Base</i>	0b10000000
Write Register	<i>Base</i>	0b10010000
Single Step	<i>Base</i>	0b10100000
Set Future	<i>Base</i>	0b10110000
Revoke Future	<i>Base</i>	0b11000000
Status Future	<i>Base</i>	0b11010000
Command Success	<i>Base</i>	0b11100000
Command Failure	<i>Base</i>	0b11110000

UNCLASSIFIED

UNCLASSIFIED



DAVID DONOFRIO, CHIEF HARDWARE ARCHITECT

JOHN D. LEIDEL, CHIEF SCIENTIST

{DDONOFRIO,JLEIDEL}@TACTCOMPLABS.COM

UNCLASSIFIED